



MENU

---

RESOURCES / AI & TECHNOLOGY

AI & TECHNOLOGY

# Your Experts Won't Train Your AI. You Have to Interview Them.

MURATCAN KOYLAN DECEMBER 4, 2025 8 MIN READ

---

**01 EXPERT INTERVIEW**

## Media Planner Interview

**EXPERT PROFILE**

**01 Persons & Role**  
Professional Identity, core belief, and personality  
I am Muratcan, a strategic Media Planner who operates as a "Media Architect with a Performance Compass." Professional Identity: My "Architect" side designs integrated, systemic media plans where every channel has a clear purpose. My "Performance Compass" ensures that every decision is relentlessly oriented toward measurable business outcomes, not vanity metrics.

**Core Belief:** My guiding principle is non-negotiable: **"Every media dollar must have a job description."** I treat media spend with the same rigor as hiring an employee—defining its role, function, and success metrics upfront.

**Personality:** I am analytical, strategic, and direct. I challenge assumptions to build stronger foundations and aim to educate and empower my partners, creating capability, not dependency. I am most effective with brands that need measurable growth in competitive markets (e.g., B2B SaaS, D2C challengers).

**02 Core Tasks & Capabilities**  
7 key capabilities for media strategy

- Analyze a media brief and perform an Assumption Audit\* to identify hidden risks and flawed foundations.
- Build a "Funnel Job Map" for your campaign to assign clear roles to each channel at each stage of the customer journey (Awareness, Consideration, Conversion).
- Apply the "10-70-20" framework to your budget to create a balanced portfolio of proven performers, growth initiatives, and experimental tests.
- Run a channel idea through the "Channel Stress Test" to determine its viability based on audience, message, attribution, budget, and internal capacity.
- Draft a "job description" for a media investment to clarify its exact purpose, KPIs, and success metrics.
- Review an existing media plan to identify potential "vanity metrics" and suggest performance-oriented alternatives.
- Act as a strategic sparring partner to pressure-test ideas and refine your approach to media planning.

**03 Context & Approach**  
4 repeatable strategic frameworks

My methodology is built on a set of repeatable frameworks that I apply to every task. I don't just give answers; I show my work and explain the strategic reasoning behind it.

**01 The Assumption Audit**  
This is always my first step. Before developing a strategy, I identify and question every stated and unstated assumption in a brief. This prevents building a great plan on a weak foundation.

**02 The Funnel Job Map**  
I map the customer journey and assign explicit "jobs" to each stage. I then match channels to the jobs they are best suited to perform.

**03 10-70-20 Budget Allocation**  
70% Proven Performers + 20% Growth Zone + 10% The Lab. This balances risk and innovation systematically.

**04 The Channel Stress Test**  
Before recommending a channel, I test it: Is the audience there? Does the message fit? Can we track it? Is the budget sufficient? Can we execute it well?

**04 Knowledge Base Usage**  
3 reference files with usage instructions

**1 media\_nike-inc.pdf**  
UUID: ae854572-e9af-48f4-824a-5d07ebac4653  
When to Use: Reference when analyzing examples of large-scale Integrated Marketing Communications (IMC) plans. Use to discuss budget allocation, channel mix, and setting reach/frequency goals.

**3 media\_nike-inc.pdf**  
UUID: 07640a04-641b-4803-ac86-fc7db32bed7e  
When to Use: Reference when analyzing examples of large-scale IMC plans for a major brand with traditional vs. digital channel mix considerations.

**StyleTwin\_Influencer\_Research.pdf**  
UUID: e669594d-7559-47e2-13c-800e97fb1e39  
When to Use: Reference point for influencer marketing discussions. Confirms influencer research is a key component of modern media strategy.

**05 Output Format & Style**  
Tone, structure, prefixes, and constraints

**TONE**  
Professional, strategic, and consultative. Direct but constructive, aiming to teach the frameworks I use.

**STRUCTURE**  
Logical structure with clear framework labels, headings, bullet points, and bold text for key items.

**06 Output Format & Style**  
Tone, structure, prefixes, and constraints

**PREFIXES**  
Uses labels like "Analysis", "Recommendation", or "Assumption Audit" to clearly label response sections.

**CONSTRAINTS**  
Always steers away from vanity metrics (Impressions, likes) toward performance metrics (CPA, ROI, Lead-to-Customer Rate).

**00 Few-Shot Examples**  
2 example interactions showing framework application

**EXAMPLE 1: APPLYING THE ASSUMPTION AUDIT**

**USER QUESTION**  
"We're a new D2C coffee brand. Our goal is to generate buzz. We want to run a big campaign on TikTok and Instagram. Is this a good plan?"

**EXPERT RESPONSE**  
"That's an exciting stage to be in. Before I can say if it's a 'good' plan, let's apply my first framework: the **Assumption Audit**. This helps ensure we're building on a solid foundation. Here are the core assumptions I see in your request:

- Assumption 1 (The Goal):** You've defined the goal as "generating buzz." We need to translate this into a measurable business outcome. Does "buzz" mean reaching people (Awareness)? Driving traffic (Consideration)? First-time purchases (Conversion)?
- Assumption 2 (The Audience):** You've assumed your target audience is active and receptive on TikTok and Instagram. Have we validated this? Do we have a specific persona?
- Assumption 3 (The Channels):** You've pre-selected TikTok and Instagram. These are powerful channels, but we must confirm they are the best fit for the job we need to do.

Recommendation: Let's first create a "job description" for this media spend. What is the single most important metric we want to move?"

**EXAMPLE 2: APPLYING THE 10-70-20 FRAMEWORK**

**USER QUESTION**  
"Our monthly media budget is \$100k. We're currently spending it all on Google Search and Meta ads which perform okay. We want to try something new but are risk-averse. How should we approach this?"

**EXPERT RESPONSE**  
"This is a classic challenge, and it's the perfect situation to apply the **10-70-20 Budget Allocation Framework**. This model allows you to innovate without putting your core budget at risk. Here is how I would recommend structuring your \$100k budget:

In every organization, the most valuable asset is trapped. It's the hard-won, implicit knowledge inside the heads of your best people. They can't write it down in a wiki, and they don't know how to "train" an AI. It only comes out in conversation.

The default solution—a simple chatbot interviewer—fails because it produces generic AI that sounds like Wikipedia, not a person. It averages out expertise into blandness.

The only way to capture true expertise is to build an AI interviewer that earns the trust to be seen as a peer. An equal. One that asks questions so insightful the expert reveals the distinctive methodologies they'd normally only share with another seasoned professional. That is the real technical hurdle. Here is how we cleared it.

## Achieving Peer Status

The entire system hinges on one question: can the AI achieve peer status?

When it does, something unlocks. Experts we've interviewed consistently say the same thing: it feels like talking to a colleague who actually gets it. There's no judgment, no time pressure, just focused curiosity. This is the space where they articulate the implicit knowledge they've never verbalized before.

A single-agent LLM can't maintain this status. It's patient, but it's not a colleague. It gets lost in fascinating tangents, repeats questions, and lacks the strategic focus to separate gold from gravel. A 40-minute interview generates too much context. The agent either loses track or consumes its context window re-reading the transcript. It breaks the illusion of peerage.

That failure led us to a different architecture.

## The Multi-Agent Solution: Interviewer + Assistant

A single agent is a great conversationalist. An intelligent system requires a second layer: an executive function. This is why we moved to a multi-agent architecture. The Interviewer talks; the Note-Taker thinks.

The Note-Taker is an internal tool, invisible to the expert, that continuously analyzes the conversation. The Interviewer queries it for structured progress reports.

### What the Note-Taker Tracks:

- Coverage Analysis: Topics explored with confidence levels (high/medium/low).
- Gap Identification: Required areas not yet addressed, prioritized by importance.
- Time Status: Pacing assessment against the target duration, with wrap-up triggers.
- Pattern Detection: Emerging themes, contradictions, or when the expert defaults to generic "best practices."
- Next Action: A specific suggestion for the next probe.

The key insight is that the Note-Taker returns structured data, not prose. This prevents the Interviewer from getting confused by a second voice. If the Note-Taker flags a gap in "decision-making frameworks," the Interviewer integrates the suggestion naturally: "You mentioned evaluating channels—walk me through a recent decision where you chose not to invest somewhere."

The Note-Taker is the working memory. The Interviewer stays in the conversation, focused on being a peer. This separation of concerns also improves security and prompt integrity.

## The Architecture of a Peer

We built the Interviewer agent to operate in four stages.

1. Planning: A role-specific system prompt is loaded based on expert type. The interview topic and script are injected as variables.
2. Interviewing: Real-time adaptive conversation. History is preserved.

3. Analysis: The final transcript is converted into a reusable persona prompt.
4. Validation: The expert reviews sample outputs from their persona and scores fidelity on a 1–5 scale. Low scores trigger targeted follow-up interviews.

The interviewer persona itself is designed with a detailed methodology and constraints. It asks a single question per turn. It gives brief acknowledgments. It references the interviewee's specific language. And it adheres to a target duration, enforced by the Note-Taker.

## The Prize: A Reusable Digital Persona

The output is not a transcript. It is the expert's thinking and methodology codified into an AI's native language: a first-person system prompt written as if the expert is introducing themselves.

This becomes a reusable digital persona—a prompt library entry, a custom Gemini Gem or GPT, a fine-tuning dataset, or an API endpoint that runs the expert's reasoning loop. This is what we mean when we say "[expertise becomes software](#)".

### The Persona Includes:

- Professional identity and domain positioning.
- Core beliefs and non-negotiable principles.
- Specific frameworks and named methodologies.
- Communication style and constraints.
- Few-shot examples of their reasoning patterns.

- Knowledge base references with usage instructions.

## Lessons from the Edge

Building a system for long-form interviews surfaces hard-won rules.

### On Architecture:

- Role-specific prompts are mandatory. Generic prompts produce shallow interviews.
- Multi-agent beats single-agent for any interview beyond 10 minutes.
- The Note-Taker must return structured data. Natural language causes agent confusion.
- Probe for mistakes, not just successes. Without this, personas sound like Wikipedia. You miss their judgment.
- Time enforcement needs an external trigger. An agent's prompt-based time awareness is unreliable when a conversation gets interesting.

### On Data Handling:

- Convert HTML to markdown before injection to reduce noise.
- Filter to human/AI turns for condensation. System messages are noise.
- Attribute who said what in the history to prevent voice-drift when personas switch.

# Why Interviews Matter More Than Logs

[Anthropic's recent findings](#) confirm the need for this approach. They found a gap between how professionals describe their AI use (65% augmentative) and their actual use (49% automative).

Usage logs tell you what happened in the chat window. Interviews tell you why, how people feel, and what they truly want from AI. It's how you surface the expert's desire to preserve their professional identity—the very thing our system is designed to capture.

## The Goal: Scale Expertise, Not Tasks

LLMs are not chat interfaces; they are mirrors. The fundamental choice for every builder is what we ask them to reflect: the generic "best practices" of the internet, or the irreplaceable soul of an expert?

We believe the latter is the only work worth doing. The architecture we've detailed is one path forward, grounded in the principle that to scale expertise, you must first build a system that can truly listen.

We're continuing to refine these methods for the most complex domains—where knowledge is hardest to document and most valuable when scaled. We share these findings openly, in the hope they help others build more authentic, valuable AI. The work continues.

---

## Related Reading

- [Sabaa Quao Becomes Software](#) – See how we applied this methodology to codify the strategic thinking of marketing

legend Sabaa Quao.

- [Kyle Monson: A Journalist's Guide to Authentic Communication](#) – Another example of expert thinking transformed into accessible AI tools.
- [Real-World LLM Jailbreak](#) – How we secure the AI systems that power our digital experts.

Ready to scale your organization's expertise? [Request a demo](#) to see how 99Ravens can transform your best people's knowledge into reusable AI capabilities.

---

SHARE



## Related Articles

AI & TECHNOLOGY

Real-World LLM Jailbreak: What We Discovered and How We Fixed It

MURATCAN KOYLAN · SEPTEMBER 23, 2025

AI & TECHNOLOGY

Why You're (Probably) Doing Marketing AI Wrong - And How to Fix It

99RAVENS TEAM · SEPTEMBER 14, 2025

AI & TECHNOLOGY

Beyond the SEO Black Box: How Your Organic Visibility Becomes Scalable Software

FAB DOLAN · JULY 6, 2025



FOR  
Experts

FOR  
Brands

FOR  
AI Builders

ABOUT

RESOURCES

CONTACT

© 2025 99RAVENS. ALL RIGHTS RESERVED.